

Logistic Regression and Neural Networks


COSC 757: Spring 2016

Logistic Regression

- Linear regression not appropriate where response variable is categorical
- Alternatively, Logistic Regression method describes relationship between categorical response and set of predictors
- Specifically, we explore applications with dichotomous response
- **Example: Suppose researchers interested in potential relationship between patient age and presence/absence of disease**
- Data set includes 20 patients

Non-Linear Relationships

- Plot shows least squares regression line (straight) and logistic regression line (curved) for *disease* on *age*
- Least squares, assumes linear relationship between variables
- In contrast, logistic regression line assumes non-linear relationship between predictor and response
- Patient 11 estimation errors (vertical lines) shown
- Patient 11's estimation error greater for linear regression versus logistic regression
- Thus, for this point, and many others, logistic regression does a better job of estimating *disease*



Sigmoid Function

- How is logistic regression line derived?
- First, $E(Y|x)$ is the conditional mean of Y , given x
- Equals expected value of response, for given predictor value
- Recall linear regression, where response random variable defined as:

$$Y = \beta_0 + \beta_1 x + \varepsilon$$

- Since ε has mean = 0, the conditional mean of Y $E(Y|x)$ equals:

$$E(Y | x) = \beta_0 + \beta_1 x$$

Sigmoid Function (continued)

- Denote $E(Y|x)$ as $\pi(x)$, where conditional mean for logistic regression takes form:

$$\pi(x) = \frac{e^{\beta_0 + \beta_1 x}}{1 + e^{\beta_0 + \beta_1 x}}$$

- Function forms s-shaped (sigmoidal) curves, which are non-linear
- Logistic function models dichotomous data well, because of simplicity and interpretability

Logistic Regression and Error

- $\pi(x)$ interpreted as probability *disease* (positive outcome) present for records $X = x$
- Furthermore, $1 - \pi(x)$ interpreted as probability *disease* (positive outcome) not present
- Recall linear regression error term ε normally distributed, with mean = 0 and constant variance
- However, assumptions regarding error term different for logistic regression

Logistic Regression and Error

- Because response dichotomous, errors take one of two forms:
- $Y = 1$ (disease present)
 - Occurs with probability $\pi(x)$, probability response positive
 - $\varepsilon = 1 - \pi(x)$ represents vertical distance between point $Y = 1$ and curve $\pi(x)$ below, for $X = x$
- $Y = 0$ (disease not present)
 - Occurs with probability $1 - \pi(x)$, probability response negative
 - $\varepsilon = 0 - \pi(x) = -\pi(x)$, which represents vertical distance between point $Y = 0$ and curve $\pi(x)$ above, for $X = x$

Logit Transformation

- Variance of $\varepsilon = \pi(x) \cdot (1 - \pi(x))$, variance of binomial distribution
- Therefore, logistic regression response $Y = \pi(x) + \varepsilon$ assumed to follow binomial distribution with probability success = $\pi(x)$
- Transformation for logistic regression, *logit transformation*, defined as:

$$g(x) = \ln \left[\frac{\pi(x)}{1 - \pi(x)} \right] = \beta_0 + \beta_1 x$$

- Includes useful properties for linearity, continuity, and ranges from positive to negative infinity

Maximum Likelihood Estimation

- Linear regression has closed form solution
- However, closed form solution does not exist for estimating logistic regression coefficients
- Therefore, maximum likelihood estimation finds parameter estimates
- Likelihood function is function of β_i parameters, which expresses probability of observed data, x

$$l(\boldsymbol{\beta} | x), \text{ where } \boldsymbol{\beta} = \beta_0, \beta_1, \dots, \beta_m$$

- Maximum likelihood estimators determined by finding values for β_i , which maximize likelihood function
- These are parameters most likely favored by data

Maximum Likelihood Estimation

- Probability of positive response, given data:
 $\pi(x) = P(Y=1 \mid x)$
- Probability of negative response, given data:
 $1 - \pi(x) = P(Y=0 \mid x)$
- Now, observations where response positive ($X_i=x_i, Y_i=1$) contribute $\pi(x)$ to likelihood, while those with negative response ($X_i=x_i, Y_i=0$) contribute $1 - \pi(x)$
- Since observations assumed independent and take on values $Y_i = 0$ or 1 , likelihood function expressed as product of terms:

$$l(\boldsymbol{\beta} \mid x) = \prod_{i=1}^n [\pi(x_i)]^{y_i} [1 - \pi(x_i)]^{1-y_i}$$

Log-likelihood Form

- Log-likelihood form $L(\beta \mid x)$ more computationally tractable:

$$L(\beta \mid x) = \ln [l(\beta \mid x)] = \sum_{i=1}^n \{y_i \ln[\pi(x_i)] + (1 - y_i) \ln[1 - \pi(x_i)]\}$$

- Finally, maximum likelihood estimates found by differentiating $L(\beta \mid x)$ with respect to each parameter β_i , and setting result = 0
- Iterative weighted least squares method applied, since closed form solutions for differentiations non-existent

Logistic Regression Example

Logistic Regression Table

Predictor	Coef	StDev	Z	P	Odds	95% CI	
					Ratio	Lower	Upper
Constant	-4.372	1.966	-2.22	0.026			
Age	0.06696	0.03223	2.08	0.038	1.07	1.00	1.14

Log-Likelihood = -10.101

Test that all slopes are zero: G = 5.696, DF = 1, P-Value = 0.017

- Coefficients (maximum likelihood estimates) of unknown parameters β_0 and β_1 , given as
 - $b_0 = -4.372$ and
 - $b_1 = 0.06696$, respectively

Logistic Regression Example

Logistic Regression Table

Predictor	Coef	StDev	Z	P	Odds	95% CI	
					Ratio	Lower	Upper
Constant	-4.372	1.966	-2.22	0.026			
Age	0.06696	0.03223	2.08	0.038	1.07	1.00	1.14

Log-Likelihood = -10.101

Test that all slopes are zero: G = 5.696, DF = 1, P-Value = 0.017

- Here, $\pi(x)$ estimated as:

$$\hat{\pi}(x) = \frac{e^{\hat{g}(x)}}{1 + e^{\hat{g}(x)}} = \frac{e^{-4.372 + 0.06696 (age)}}{1 + e^{-4.372 + 0.06696 (age)}}$$

- With estimated logit:

$$\hat{g}(x) = -4.372 + 0.06696(age)$$

Logistic Regression Example

Logistic Regression Table

Predictor	Coef	StDev	Z	P	Odds Ratio	95% CI	
						Lower	Upper
Constant	-4.372	1.966	-2.22	0.026			
Age	0.06696	0.03223	2.08	0.038	1.07	1.00	1.14

Log-Likelihood = -10.101

Test that all slopes are zero: G = 5.696, DF = 1, P-Value = 0.017

- Using these equations, estimated probability disease present in patient, given their age derived
- Example: Estimate probability disease present in particular patient, age = 50

$$\hat{g}(x) = -4.372 + 0.06696(50) = -1.024 \quad \hat{\pi}(x) = \frac{e^{\hat{g}(x)}}{1 + e^{\hat{g}(x)}} = \frac{e^{-1.024}}{1 + e^{-1.024}} = 0.26$$


- Estimated probability 50-year old patient has disease = 26%, with probability disease not present = 74%

More on Logistic Regression

- Logistic Regression Tutorial in R
- <http://ww2.coastal.edu/kingw/statistics/R-tutorials/logistic.html>


Neural Networks

- Complex learning systems recognized in animal brains
- Single neuron has a simple structure
- Interconnected sets of neurons perform complex learning tasks
- Artificial Neural Networks attempt to replicate non-linear learning found in nature




Before Neural Networks: The Perceptron

- Linear classifier
- Find decision boundary between classes
- Steps:
 - Initialize $w = 0$
 - For each x predict positive iff $w_t \cdot x > 0$
 - On mistake update w
 - Positive: $w_{t+1} \leftarrow w_t + x$
 - Negative: $w_{t+1} \leftarrow w_t - x$
 - $t \leftarrow t + 1$



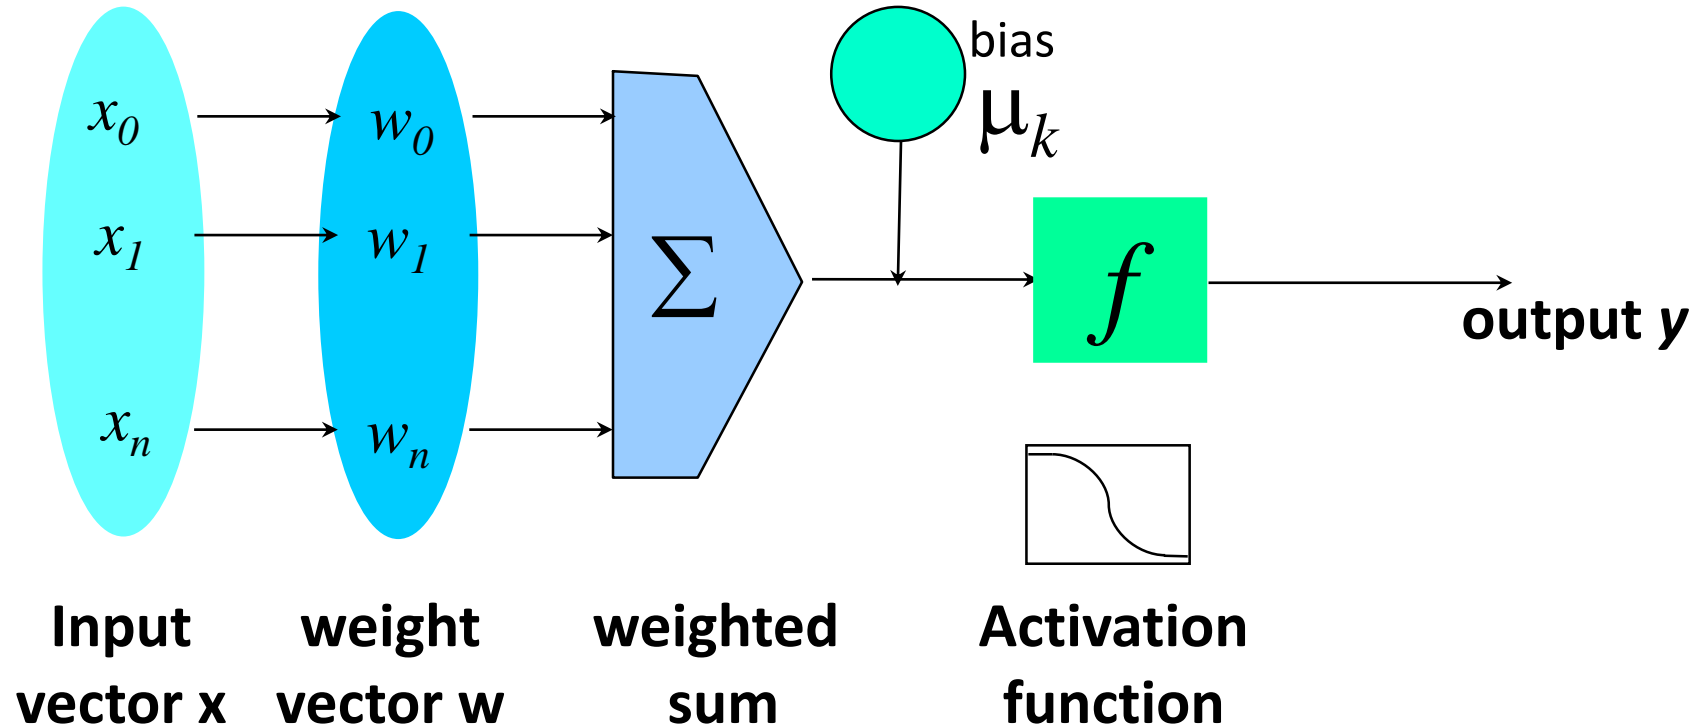
$$f(x) = \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$



Classification by Backpropagation

- Backpropagation: A **neural network** learning algorithm
- Started by psychologists and neurobiologists to develop and test computational analogues of neurons
- A neural network: A set of connected input/output units where each connection has a **weight** associated with it
- During the learning phase, the **network learns by adjusting the weights** so as to be able to predict the correct class label of the input tuples
- Also referred to as **connectionist learning** due to the connections between units

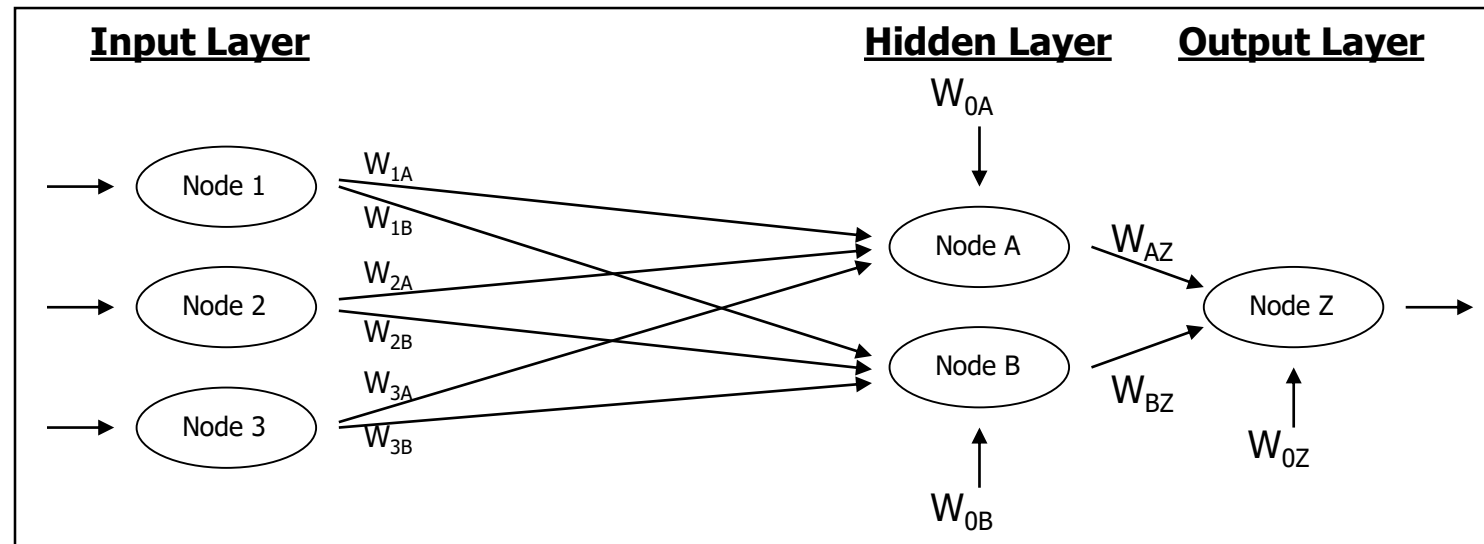
Neuron: A Hidden/Output Layer Unit



- An n -dimensional input vector \mathbf{x} is mapped into variable y by means of the scalar product and a nonlinear function mapping
- The inputs to unit are outputs from the previous layer. They are multiplied by their corresponding weights to form a weighted sum, which is added to the bias associated with unit. Then a nonlinear activation function is applied to it.

Neural Network Components

- Neural Network consists of layered, feedforward, completely connected network of nodes
- Feedforward restricts network flow to single direction
- Flow does not loop or cycle
- Network composed of two or more layers



Neural Network Components

- Most networks have Input, Hidden, Output layers
- Network may contain more than one hidden layer
- Network is completely connected
- Each node in given layer, connected to every node in next layer
- Every connection has weight (W_{ij}) associated with it
- Weight values randomly assigned 0 to 1 by algorithm
- Number of input nodes dependent on number of predictors
- Number of hidden and output nodes configurable

Defining a Network Topology

- Decide the **network topology**: Specify # of units in the *input layer*, # of *hidden layers* (if > 1), # of units in *each hidden layer*, and # of units in the *output layer*
- Normalize the input values for each attribute measured in the training tuples to $[0.0—1.0]$
- One **input** unit per domain value, each initialized to 0
- **Output**, if for classification and more than two classes, one output unit per class is used
- Once a network has been trained and its accuracy is **unacceptable**, repeat the training process with a *different network topology* or a *different set of initial weights*

Neural Network Layers

- Hidden Layer
 - How many nodes in hidden layer?
 - Large number of nodes increases complexity of model
 - Detailed patterns uncovered in data
 - Leads to overfitting, at expense of generalizability
 - Reduce number of hidden nodes when overfitting occurs
 - Increase number of hidden nodes when training accuracy unacceptably low
- Input Layer
 - Input layer accepts values from input variables
 - Values passed to hidden layer nodes
 - Input layer nodes lack detailed structure compared to hidden and output layer nodes

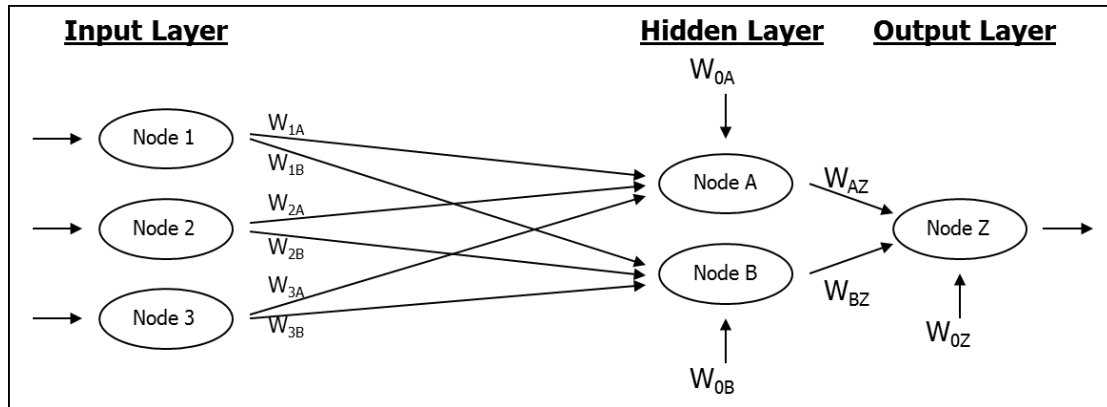
Nodes and Weights

- Combination function produces linear combination of node inputs and connection weights to single scalar value
- For a given node j :

$$\text{net}_j = \sum_i W_{ij} x_{ij} = W_{0j} x_{0j} + W_{1j} x_{1j} + \dots + W_{Ij} x_{Ij}$$

- where
 - x_{ij} is i th input to node j
 - W_{ij} is weight associated with i th input node
 - and there are $I + 1$ inputs to node j
 - x_1, x_2, \dots, x_I are inputs from upstream nodes
 - x_0 is constant input value = 1.0
 - Each input node has extra input $W_{0j}x_{0j} = W_{0j}$

Nodes and Weights



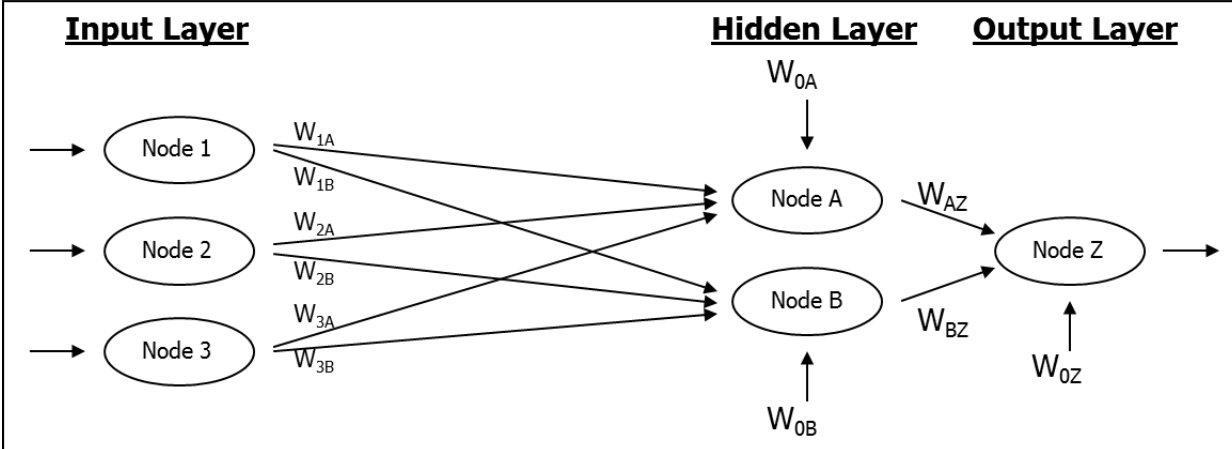
$x_0 = 1.0$	$W_{0A} = 0.5$	$W_{0B} = 0.7$	$W_{0Z} = 0.5$
$x_1 = 0.4$	$W_{1A} = 0.6$	$W_{1B} = 0.9$	$W_{AZ} = 0.9$
$x_2 = 0.2$	$W_{2A} = 0.8$	$W_{2B} = 0.8$	$W_{BZ} = 0.9$
$x_3 = 0.7$	$W_{3A} = 0.6$	$W_{3B} = 0.4$	

- The scalar value computed for hidden layer Node A equals

$$\text{net}_A = \sum_i W_{iA} x_{iA} = W_{0A}(1.0) + W_{1A}x_{1A} + W_{2A}x_{2A} + W_{3A}x_{3A} =$$

$$0.5 + 0.6(0.4) + 0.8(0.2) + 0.6(0.7) = 1.32$$

- For Node A, $\text{net}_A = 1.32$ is the input to activation function
- This activation is analogous to how neurons “fire” nonlinearly in biological organisms




$x_0 = 1.0$	$W_{0A} = 0.5$	$W_{0B} = 0.7$	$W_{0Z} = 0.5$
$x_1 = 0.4$	$W_{1A} = 0.6$	$W_{1B} = 0.9$	$W_{AZ} = 0.9$
$x_2 = 0.2$	$W_{2A} = 0.8$	$W_{2B} = 0.8$	$W_{BZ} = 0.9$
$x_3 = 0.7$	$W_{3A} = 0.6$	$W_{3B} = 0.4$	

- Firing response not necessarily linearly related to increase in input stimulation
- Artificial Neural Networks model behavior using non-linear activation function
- Sigmoid function most commonly used

$$y = \frac{1}{1 + e^{-x}}$$

- In Node A, activation function takes netA = 1.32 as input and produces output

$$y = \frac{1}{1 + e^{-1.32}} = 0.7892$$



$x_0 = 1.0$	$W_{0A} = 0.5$	$W_{0B} = 0.7$	$W_{0Z} = 0.5$
$x_1 = 0.4$	$W_{1A} = 0.6$	$W_{1B} = 0.9$	$W_{AZ} = 0.9$
$x_2 = 0.2$	$W_{2A} = 0.8$	$W_{2B} = 0.8$	$W_{BZ} = 0.9$
$x_3 = 0.7$	$W_{3A} = 0.6$	$W_{3B} = 0.4$	


- Node A outputs 0.7892 along connection to Node B, and becomes component of netZ
- Before netZ is computed, contribution from Node B required

$$\text{net}_B = \sum_i W_{iB} x_{iB} = W_{0B}(1.0) + W_{1B}x_{1B} + W_{2B}x_{2B} + W_{3B}x_{3B} = 0.7 + 0.9(0.4) + 0.8(0.2) + 0.4(0.7) = 1.5$$

- then

$$f(\text{net}_B) = \frac{1}{1 + e^{-1.5}} = 0.8176$$

- Node Z combines outputs from Node A and Node B, through netZ, a weighted sum, using weight associated to the connections between nodes



$x_0 = 1.0$	$W_{0A} = 0.5$	$W_{0B} = 0.7$	$W_{0Z} = 0.5$
$x_1 = 0.4$	$W_{1A} = 0.6$	$W_{1B} = 0.9$	$W_{AZ} = 0.9$
$x_2 = 0.2$	$W_{2A} = 0.8$	$W_{2B} = 0.8$	$W_{BZ} = 0.9$
$x_3 = 0.7$	$W_{3A} = 0.6$	$W_{3B} = 0.4$	

- Inputs to Node Z not data attribute values
- Rather, outputs are from sigmoid function in upstream nodes


$$\text{net}_Z = \sum_i W_{iZ} x_{iZ} = W_{0Z} (1.0) + W_{AZ} x_{AZ} + W_{BZ} x_{BZ} = 0.5 + 0.9(0.7892) + 0.9(0.8176) = 1.9461$$

- then

$$f(\text{net}_Z) = \frac{1}{1 + e^{-1.9461}} = 0.8750$$

- Value 0.8750 output from Neural Network on first pass
- Represents predicted value for target variable, given first observation

Sigmoid Function



- Sigmoid function combines nearly linear, curvilinear, and nearly constant behavior depending on input value
- Function nearly linear for domain values $-1 < x < 1$
- Becomes curvilinear as values move away from center
- At extreme values, $f(x)$ is nearly constant
- Moderate increments in x produce variable increase in $f(x)$, depending on location of x
- Sometimes called “Squashing Function”
- Takes real-valued input and returns values $[0, 1]$

Back Propagation

- Neural Networks are supervised learning method
- Require target variable
- Each observation passed through network results in output value
- **Output value compared to actual value of target variable**
- $(\text{Actual} - \text{Output}) = \text{Error}$
- Prediction error analogous to residuals in regression models
- Most networks use Sum of Squares (SSE) to measure how well predictions fit target values

$$SSE = \sum_{\text{Records}} \sum_{\text{OutputNodes}} (\text{actual} - \text{output})^2$$

Back Propagation


- Squared prediction errors summed over all output nodes, and all records in data set
- Model weights constructed that minimize SSE
- Actual values that minimize SSE are unknown
- Weights estimated, given the data set
- Unlike least-squares regression, no closed-form solution exists for minimizing SSE

Gradient Descent

- Gradient Descent Method determines set of weights that minimize SSE
- Given a set of m weights $w = w_1, w_2, \dots, w_m$ in network model
- Find values for weights that, together, minimize SSE
- Gradient Descent determines direction to adjust weights, that decreases SSE
- Gradient of SSE, with respect to vector of weights w , is vector derivative:


$$\nabla SSE(w) = \left[\frac{\partial SSE}{\partial w_0}, \frac{\partial SSE}{\partial w_1}, \dots, \frac{\partial SSE}{\partial w_m} \right]$$

Gradient Descent



Gradient Descent

- Basic Intuition:
 - Calculate slope at current position
 - If slope is negative, move right
 - If slope is positive, move left
 - (Repeat until slope == 0)



Gradient Descent

- Develop rule defining movement from current w_1 to optimal value w^*1

$$w_{NEW} = w_{CURRENT} + \Delta w_{CURRENT}$$

where

$\Delta w_{CURRENT}$ is change in current location w

- If current weight near w_{1L} , increasing w approaches w^*1
- If current weight near w_{1R} , decreasing w approaches w^*1
- Gradient of SSE, with respect to weight $w_{CURRENT}$, is slope of SSE curve at $w_{CURRENT}$
- Value $w_{CURRENT}$ close to w_{1L} , slope is negative
- Value $w_{CURRENT}$ close to w_{1R} , slope is positive

Gradient Descent

- **Direction** for adjusting $w_{CURRENT}$ is negative sign of derivative at SSE at $w_{CURRENT}$

$$- \text{sign}\left(\frac{\partial SSE}{\partial w_{CURRENT}}\right)$$

- To adjust, use **magnitude** of the derivative of SSE at $w_{CURRENT}$
- **When curve steep, adjustment is large**
- **When curve nearly flat, adjustment is small**

$$\Delta w_{CURRENT} = -\eta \left(\frac{\partial SSE}{\partial w_{CURRENT}} \right)$$

- Learning Rate (Greek “eta”) has values $[0, 1]$

Back Propagation Rules

- Back-propagation percolates prediction error for record back through network
- Partitioned responsibility for prediction error assigned to various connections
- Weights of connections adjusted to decrease error, using Gradient Descent Method

$$w_{ij,NEW} = w_{ij,CURRENT} + \Delta w_{ij}$$

where

$$\Delta w_{ij} = \eta \delta_j x_{ij}$$

η = learning rate

x_{ij} = signifies i th input to node j

δ_j = represents responsibility for a particular error belonging to node j

Back Propagation Rules

- Error responsibility computed using partial derivative of the sigmoid function with respect to net_j

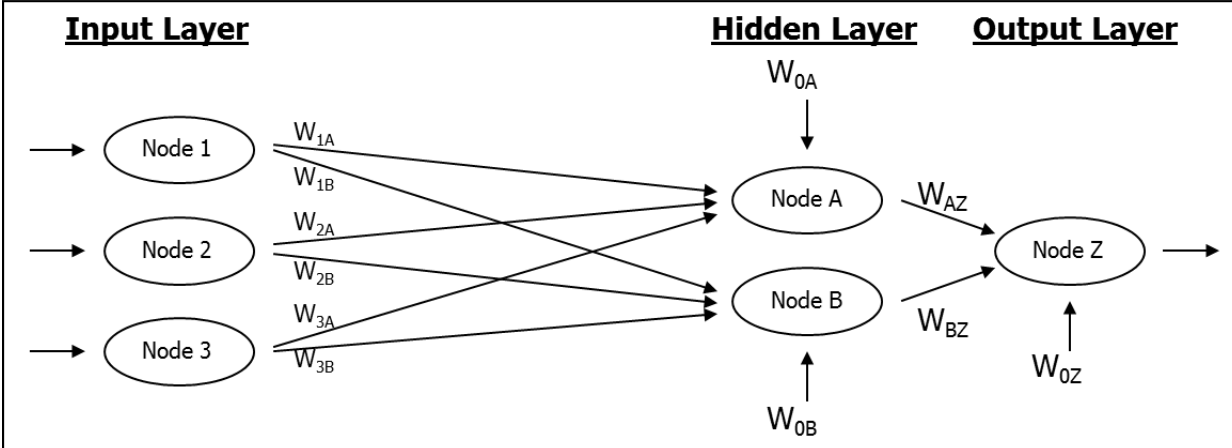
- Values take one of two forms

$$\delta_j = \begin{cases} \text{output}_j(1 - \text{output}_j)(\text{actual}_j - \text{output}_j) & \text{for output layer nodes} \\ \text{output}_j(1 - \text{output}_j) \sum_{\text{DOWNSTREAM}} W_{jk} \delta_j & \text{for hidden layer nodes} \end{cases}$$

- where

$$\sum_{\text{DOWNSTREAM}} W_{jk} \delta_j$$


- refers to the weighted sum of error responsibilities for nodes downstream
- Rules show why input values require normalization
- Large input values x_{ij} would dominate weight adjustment
- Error propagation would be overwhelmed, and learning stifled



$x_0 = 1.0$	$W_{0A} = 0.5$	$W_{0B} = 0.7$	$W_{0Z} = 0.5$
$x_1 = 0.4$	$W_{1A} = 0.6$	$W_{1B} = 0.9$	$W_{AZ} = 0.9$
$x_2 = 0.2$	$W_{2A} = 0.8$	$W_{2B} = 0.8$	$W_{BZ} = 0.9$
$x_3 = 0.7$	$W_{3A} = 0.6$	$W_{3B} = 0.4$	

- Recall that first pass through network yielded output = 0.8750
- Assume actual target value = 0.8, and learning rate, $\eta = 0.01$
- Prediction error = $0.8 - 0.8750 = -0.075$
- Neural Networks use stochastic (or online) back-propagation
- Weights updated after each record processed by network
- Error responsibility for Node Z, an output node, found first

$$\delta_Z = \text{output}_Z (1 - \text{output}_Z) (\text{actual}_Z - \text{output}_Z) = 0.875(1 - 0.875)(0.8 - 0.875) = -0.0082$$



$x_0 = 1.0$	$W_{0A} = 0.5$	$W_{0B} = 0.7$	$W_{0Z} = 0.5$
$x_1 = 0.4$	$W_{1A} = 0.6$	$W_{1B} = 0.9$	$W_{AZ} = 0.9$
$x_2 = 0.2$	$W_{2A} = 0.8$	$W_{2B} = 0.8$	$W_{BZ} = 0.9$
$x_3 = 0.7$	$W_{3A} = 0.6$	$W_{3B} = 0.4$	


- Now adjust “constant” weight w_{0Z} using rules

$$\Delta W_{0Z} = \eta \delta_Z (1) = 0.1(-0.0082)(1) = -0.00082$$

$$W_{0Z,NEW} = W_{0Z,CURRENT} + \Delta W_{0Z} = 0.5 - 0.00082 = 0.49918$$

- Move upstream to Node A, a hidden layer node

$$\begin{aligned} \delta_A &= \text{output}_A (1 - \text{output}_A) \sum_{DOWNSTREAM} W_{jk} \delta_j \\ &= 0.7892(1 - 0.7892)(0.9)(-0.0082) = -0.00123 \end{aligned}$$



$x_0 = 1.0$	$W_{0A} = 0.5$	$W_{0B} = 0.7$	$W_{0Z} = 0.5$
$x_1 = 0.4$	$W_{1A} = 0.6$	$W_{1B} = 0.9$	$W_{AZ} = 0.9$
$x_2 = 0.2$	$W_{2A} = 0.8$	$W_{2B} = 0.8$	$W_{BZ} = 0.9$
$x_3 = 0.7$	$W_{3A} = 0.6$	$W_{3B} = 0.4$	


- Adjust weight using back-propagation rules

$$\Delta W_{AZ} = \eta \delta_Z (OUTPUT_A) = 0.1(-0.0082)(0.7892) = -0.000647$$

$$W_{AZ,NEW} = W_{AZ,CURRENT} + \Delta W_{AZ} = 0.9 - 0.000647 = 0.899353$$

- Connection weight between Node A and Node Z adjusted from 0.9 to 0.899353
- Calculate error at Node B (hidden layer node)

$$\begin{aligned} \delta_B &= output_B (1 - output_B) \sum_{DOWNSTREAM} W_{jk} \delta_j \\ &= 0.8176(1 - 0.8176)(0.9)(-0.0082) = -0.0011 \end{aligned}$$



$x_0 = 1.0$	$W_{0A} = 0.5$	$W_{0B} = 0.7$	$W_{0Z} = 0.5$
$x_1 = 0.4$	$W_{1A} = 0.6$	$W_{1B} = 0.9$	$W_{AZ} = 0.9$
$x_2 = 0.2$	$W_{2A} = 0.8$	$W_{2B} = 0.8$	$W_{BZ} = 0.9$
$x_3 = 0.7$	$W_{3A} = 0.6$	$W_{3B} = 0.4$	

- Adjust weight using back-propagation rules

$$\Delta W_{BZ} = \eta \delta_Z (OUTPUT_B) = 0.1(-0.0082)(0.8176) = -0.00067$$

$$W_{BZ,NEW} = W_{BZ,CURRENT} + \Delta W_{BZ} = 0.9 - 0.00067 = 0.89933$$

- Connection weight between Node B and Node Z adjusted from 0.9 to 0.89933
- Similarly, application of back-propagation rules continues to input layer nodes
- Weights {w1A, w2A, w3A , w0A} and {w1B, w2B, w3B , w0B} updated by process

Example Summary

- Now, all network weights in model are updated
- Each iteration based on single record from data set
- Network calculated predicted value for target variable
- Prediction error derived
- Prediction error percolated back through network
- Weights adjusted to generate smaller prediction error
- Process repeats record by record


Termination Criteria

- Many passes through data set performed before termination criterion is met
- Constantly adjusting weights to reduce prediction error
- When to terminate?
 - Stopping criterion may be computational “clock” time?
 - Short training times likely result in poor model
 - Terminate when SSE reaches threshold level?
 - Neural Networks are prone to overfitting
 - Memorizing patterns rather than generalizing

Termination Criteria

- Cross-Validation Termination Procedure
 - Retain portion of training set as “hold out” data set
 - Train network on remaining data
 - Apply weights learned from training set to validation set
 - Measure two sets of weights
 - “Current” weights for training set, “Best” weights with minimum SSE on validation set
 - Terminate algorithm when current weights has significantly greater SSE than best weights
 - However, Neural Networks not guaranteed to arrive at global minimum for SSE

Termination Criteria



- Algorithm may become stuck in local minimum
- Results in good, but not optimal solution
- Not necessarily an insuperable problem
- Multiple networks trained using different starting weights
- Best model from group chosen as “final”
- Stochastic back-propagation method acts to prevent getting stuck in local minimum
- Random element introduced to gradient descent
- Momentum term may be added to back-propagation algorithm

Learning Rate

- Recall Learning Rate (Greek “eta”) is a constant


$$0 < \eta < 1, \text{ where}$$

η = learning rate

- Helps adjust weights toward global minimum for SSE
- Small Learning Rate
 - With small learning rate, weight adjustments small
 - Network takes unacceptable time converging to solution
- Large Learning Rate
 - Suppose algorithm close to optimal solution
 - With large learning rate, network likely to “overshoot” optimal solution

Learning Rate

- w^* optimum weight for w , which has current value w_{CURRENT}
- According to Gradient Descent Rule, w_{CURRENT} adjusted in direction of w^*
- Learning rate acts as multiplier to formula $\Delta w_{\text{CURRENT}}$
- Large learning may cause w_{NEW} to jump past w^*
- w_{NEW} may be farther away from w^* than w_{CURRENT}
- Next adjusted weight value on opposite side of w^*
- Leads to oscillation between two “slopes”
- Network never settles down to minimum between them



Adjust Learning Rate as Training Process

- Learning rate initialized with large value
- Network quickly approaches general vicinity of optimal solution
- As network begins to converge, learning rate gradually reduced
- Avoids overshooting minimum

Momentum Term

- Momentum term (“alpha”) makes back-propagation more powerful, and represents inertia

$$\Delta w_{CURRENT} = -\eta \frac{\partial SSE}{\partial w_{CURRENT}} + \alpha \Delta w_{PREVIOUS}$$

- where

$\Delta w_{PREVIOUS}$ = previous weight adjustment, $0 \leq \alpha < 1$

$\alpha \Delta w_{PREVIOUS}$ = fraction of previous weight adjustment for a given weight

- Large momentum values influence $\Delta w_{CURRENT}$ to move same direction as previous adjustments
- Including momentum in back-propagation results in adjustments becoming exponential average of all previous adjustments (Reed and Marks)



$$\Delta w_{CURRENT} = -\eta \sum_{k=0}^{\infty} \alpha^k \frac{\partial SSE}{\partial w_{CURRENT-k}}, \text{ where}$$

α^k = indicates more recent adjustments exert larger influence

Momentum Term


- Large alpha values enable algorithm to remember more terms in adjustment history
- Small alpha values reduce inertial effects, and influence of previous adjustments
- With $\alpha = 0$, all previous components disappear
- Momentum term encourages adjustments in same direction
- Increases rate which algorithm approaches neighborhood of optimality
- Early adjustments in same direction
- Exponential average of adjustments, also in same direction

Momentum Term



- Weight initialized to I, with local minima at A, C
- Global minimum at B
- Small Ball
 - Small “ball”, symbolized as momentum, rolls down hill
 - Gets stuck in first trough A
 - Momentum helps find A (local minimum), but not global minimum B
- Large Ball
 - Now, large “ball” symbolizes momentum term
 - Ball rolls down curve and passes over first and second hills
 - Overshoots global minimum B because of too much momentum
 - Settles to local minimum at C
- Values of learning rate and momentum require careful consideration
- Experimentation with different values necessary to obtain best results

Why Neural Networks Work



- Neural networks have many hidden layers
- Each node is initialized in a random starting state
- This increases the ability of the neural network to cover a large search space
- Results in finding many local minima and arrive at the optimal solution.

Sensitivity Analysis

- Opacity is drawback of Neural Networks
- Flexibility enables modeling of non-linear behavior
- However, limits ability to interpret results
- No procedure exists for translating weights to decision rules
- Sensitivity Analysis measures relative influence attributes have on solution

Sensitivity Analysis Procedure

- Generate new observation xMEAN
- Each attribute value of xMEAN equals mean value for attributes in data set
- Find network output, for input xMEAN, called outputMEAN
- Attribute by attribute, vary xMEAN to reflect attribute Min and Max
- Find network output for each variation, compare to outputMEAN
- Determines attributes, varying from their Min to Max, having greater effect on network results, compared to other attributes

If you really want to delve into neural networks

- Here is a two part blog post with examples in Python!
- <http://iamtrask.github.io/2015/07/12/basic-python-network/>
- <http://iamtrask.github.io/2015/07/27/python-network-part2/>